



Figure 7 - mpdd metal_plate: multi-method per-sample comparison on the S8 common region (seed 0, K 4; region 392 x 392, 76.6% of the canvas)


metal_plate / metal_plate/test/scratches/028.png P-AP spread 0.639 over 6 method column(s) with data; selected rank 1 of 71



metal_plate / metal_plate/test/scratches/026.png P-AP spread 0.633 over 6 method column(s) with data; selected rank 2 of 71



metal_plate / metal_plate/test/scratches/027.png P-AP spread 0.622 over 6 method column(s) with data; selected rank 3 of 71



Panels: the S8 common region (the intersection of the rectangles the compared methods cover); every method map is resampled once onto that grid (linear; the GT with nearest neighbours). The number under a method label is that method's per-sample Pixel-AP on the region grid (two decimals; the JSON summary holds it in full).

Colour: magma with one shared min-max range per row across the method columns; the bar under the columns gives that row's actual range.

Selection (a display choice): per category the 3 test images whose ground truth covers at least 76 region pixels and whose per-sample Pixel-AP spread across the columns with data is largest; ties by sample id.

n/a: no per-sample dump for this unit; the column stays for row alignment and the reason is in the JSON summary.

Abbreviations: A1 J/L = controlled A1, shared/independent support rows; ADino(-rot) = AnomalyDINO, canvas frame (with rotation); PC-128/224 = PatchCore native Resize(144/256) + Crop(128/224).

Sources: baseline_common_region.csv and common_region_geometry.json (S8) plus the per-method dumps listed in the JSON summary; ground truth from the canonical caches.