




Protocol of each bar (each method keeps its own; “same protocol” here means the same measurement methodology, not a common resolution).

A1 J / A1 L B = DINOv2-B/14 @448, C = AnomalyCLIP ViT-L/14@336 @518
 canvas 448x448, no rotation, bank K; L differs from J only in the fusion order
 ADino DINOv2-S/14 @448, canvas 448x448, no rotation, bank K
 ADino-rot DINOv2-S/14 @448, canvas 448x448, 8 rotations per reference, bank Kx8
 PC-128 WideResNet50-2 layer2+3, resize 144 / imagesize 128, coreset 10%
 PC-224 WideResNet50-2 layer2+3, resize 256 / imagesize 224, coreset 10%

End to end = reference-set encoding + query-set scoring. Excluded and charged to no method: dataset construction, model loading, metrics (AP/AUROC) and disk writes; stage boundaries are anchored with torch.cuda.synchronize().

1 untimed warm-up + 3 timed repeats per unit, median reported; the spread and every raw repetition are in the JSON snapshot and no error bar is drawn.

VRAM bars: torch.cuda.max_memory_allocated (allocator), reset before each timed repeat.

Diamonds: device-wide nvidia-smi delta, larger by the CUDA context and workspaces.

Sources: 05_baselines/SPEED_VRAM_BENCH.json + .csv
 (scripts/limitation_closure_20260915/bench_inference_speed_vram.py); nothing is recomputed here. Device baseline: 2802 MB of 6144 MB before any child started.