




Figure 7 · MVTEC zipper · common-region comparison (seed 0, K 4; 392 × 392 grid; 76.6% canvas)



Panels: the S8 common region (the intersection of the rectangles the compared methods cover); every method map is resampled once onto that grid (linear; the GT with nearest neighbours). The number under a method label is that method's per-sample Pixel-AP on the region grid (two decimals; the JSON summary holds it in full).

Colour: magma with one shared min-max range per row across the method columns; the bar under the columns gives that row's actual range.

Contour (display only): each method column's score map on the region grid is quantised to 256 bins and split at the threshold that maximises the between-class variance (first maximum on ties); the contour at or above it is drawn in cyan. The identical criterion is applied to every method column, so all compared methods are treated alike. It is a display choice and is not the detection threshold behind any reported metric.

Selection (a display choice): per category the 3 test images whose ground truth covers at least 76 region pixels and whose per-sample Pixel-AP spread across the columns with data is largest; ties by sample id.

n/a: no per-sample dump for this unit; the column stays for row alignment and the reason is in the JSON summary (AnomalyDINO canvas, AnomalyDINO canvas + rotation, PatchCore native 128, PatchCore official 224 here).

Method definitions: Dual-encoder baseline, Joint matching, shared support row; Dual-encoder baseline, Independent matching, independent support rows; AnomalyDINO, canvas frame; AnomalyDINO, canvas frame plus rotation; PatchCore native Resize(144) plus CenterCrop(128); PatchCore native Resize(256) plus CenterCrop(224).

Sources: baseline_common_region.csv and common_region_geometry.json (S8) plus the per-method dumps listed in the JSON summary; ground truth from the canonical caches.