



Figure 7 · MPDD bracket_white · common-region comparison (seed 0, K 4; 392 × 392 grid; 76.6% canvas)


bracket_white / bracket_white/test/defective_painting/009.png P-AP spread 0.070 over 2 method column(s) with data; selected rank 1 of 18



bracket_white / bracket_white/test/defective_painting/000.png P-AP spread 0.044 over 2 method column(s) with data; selected rank 2 of 18



bracket_white / bracket_white/test/defective_painting/008.png P-AP spread 0.024 over 2 method column(s) with data; selected rank 3 of 18



Panels: the S8 common region (the intersection of the rectangles the compared methods cover); every method map is resampled once onto that grid (linear; the GT with nearest neighbours). The number under a method label is that method's per-sample Pixel-AP on the region grid (two decimals; the JSON summary holds it in full).

Colour: magma with one shared min-max range per row across the method columns; the bar under the columns gives that row's actual range.

Selection (a display choice): per category the 3 test images whose ground truth covers at least 76 region pixels and whose per-sample Pixel-AP spread across the columns with data is largest; ties by sample id.

n/a: no per-sample dump for this unit; the column stays for row alignment and the reason is in the JSON summary (ADino, ADino-rot, PC-128, PC-224 here).

Abbreviations: A1 J/L = controlled A1, shared/independent support rows; ADino(-rot) = AnomalyDINO, canvas frame (with rotation); PC-128/224 = PatchCore native Resize(144/256) + Crop(128/224).

Sources: baseline_common_region.csv and common_region_geometry.json (S8) plus the per-method dumps listed in the JSON summary; ground truth from the canonical caches.